

CRUDE OIL PRICE PREDICTION

November 15, 2024

1 DATA PREPROCESSING

1.1 Importing the libraries

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import tensorflow as tf
```

```
[2]: pwd
```

```
[2]: 'C:\\Users\\lenovo\\Desktop'
```

```
[3]: import pandas as pd
data= pd.read_csv(r"C:\Users\lenovo\Downloads\Crude Oil Prices Daily.csv")
print(data.head())
```

| | Date | Closing Value |
|---|----------|---------------|
| 0 | 1/2/1986 | 25.56 |
| 1 | 1/3/1986 | 26.00 |
| 2 | 1/6/1986 | 26.53 |
| 3 | 1/7/1986 | 25.85 |
| 4 | 1/8/1986 | 25.87 |

1.2 Analyze the Data

```
[4]: data.head()
```

```
[4]:
```

| | Date | Closing Value |
|---|----------|---------------|
| 0 | 1/2/1986 | 25.56 |
| 1 | 1/3/1986 | 26.00 |
| 2 | 1/6/1986 | 26.53 |
| 3 | 1/7/1986 | 25.85 |
| 4 | 1/8/1986 | 25.87 |

```
[11]: data.tail()
```

```
[11]:
```

| | Date | Closing Value |
|------|----------|---------------|
| 8218 | 7/3/2018 | 74.19 |
| 8219 | 7/4/2018 | NaN |
| 8220 | 7/5/2018 | 73.05 |
| 8221 | 7/6/2018 | 73.78 |
| 8222 | 7/9/2018 | 73.93 |

```
[12]: data.describe()
```

```
[12]:
```

| | Closing Value |
|-------|---------------|
| count | 8216.000000 |
| mean | 43.492139 |
| std | 29.616804 |
| min | 10.250000 |
| 25% | 19.577500 |
| 50% | 29.610000 |
| 75% | 63.402500 |
| max | 145.310000 |

```
data.info()
```

1.3 Handling missing values

```
[14]: data.isnull().any()
```

```
[14]: Date                False
      Closing Value       True
      dtype: bool
```

```
[15]: data.isnull().sum()
```

```
[15]: Date                0
      Closing Value       7
      dtype: int64
```

```
[16]: data.dropna(axis=0,inplace=True)
```

```
[17]: data.isnull().sum()
```

```
[17]: Date                0
      Closing Value       0
      dtype: int64
```

```
[18]: data_oil = data.reset_index()['Closing Value']
```

```
[19]: data_oil
```

```
[19]: 0      25.56
      1      26.00
      2      26.53
      3      25.85
      4      25.87
      ...
      8211    73.89
      8212    74.19
      8213    73.05
      8214    73.78
      8215    73.93
      Name: Closing Value, Length: 8216, dtype: float64
```

```
[20]: print(data_oil.isnull().sum())
      print(data_oil.shape)
```

```
0
(8216,)
```

```
[21]: data_oil.dropna(inplace=True)
      print(data_oil.isnull().sum())
      print(data_oil.shape)
```

```
0
(8216,)
```

```
[22]: print(data_oil.isnull().any())
```


```
False
```

1.4 Feature Scaling

```
[23]: from sklearn.preprocessing import MinMaxScaler
      scaler = MinMaxScaler(feature_range=(0, 1))
      data_oil = scaler.fit_transform(np.array(data_oil).reshape(-1, 1))
```

1.5 Data Visualization

```
[24]: import matplotlib.pyplot as plt # Use plt as the standard alias
      plt.title('Crude Oil Closing Prices Over Time')
      plt.plot(data_oil)
      plt.xlabel('Time')
      plt.ylabel('Closing Price')
      plt.show()
```



1.6 Splitting Data into Train and Test

```
[25]: training_size = int(len(data_oil)*0.65)
      test_size = len(data_oil) - training_size
      train_data, test_data = data_oil[0:training_size,:], data_oil[training_size:
      ↪ len(data_oil),:]
```

```
[26]: training_size, test_size
      (5340, 2876)
```

```
[26]: (5340, 2876)
```

```
[27]: train_data.shape
      (5340)
```

```
[27]: 5340
```

```
[28]: import numpy as np

      def create_dataset(dataset, time_step=1):
```

```

dataX, dataY = [], []
for i in range(len(dataset) - time_step - 1):
    a = dataset[i:(i + time_step), 0]
    dataX.append(a)
    dataY.append(dataset[i + time_step, 0])
return np.array(dataX), np.array(dataY)

```

```

[29]: time_step = 10
      X_train, y_train = create_dataset(train_data, time_step)
      X_test, ytest = create_dataset(test_data, time_step)

```

```

[30]: print(X_train.shape,y_train.shape)

```

```

(5329, 10) (5329,)

```

```

[31]: print(X_test.shape,y_train.shape)

```

```

(2865, 10) (5329,)

```

```

[32]: X_train

```

```

[32]: array([[0.11335703, 0.11661484, 0.12053902, ..., 0.10980305, 0.1089886 ,
              0.11054346],
             [0.11661484, 0.12053902, 0.11550422, ..., 0.1089886 , 0.11054346,
              0.10165852],
             [0.12053902, 0.11550422, 0.1156523 , ..., 0.11054346, 0.10165852,
              0.09906708],
             ...,
             [0.36731823, 0.35176958, 0.36080261, ..., 0.36391234, 0.37042796,
              0.37042796],
             [0.35176958, 0.36080261, 0.35354657, ..., 0.37042796, 0.37042796,
              0.37879461],
             [0.36080261, 0.35354657, 0.35295424, ..., 0.37042796, 0.37879461,
              0.37916482]])

```

```

[33]: X_train = X_train.reshape(X_train.shape[0],X_train.shape[1],1)
      X_test = X_test.reshape(X_test.shape[0],X_test.shape[1],1)

```

2 Model Building

2.1 Importing the Model Building Libraries

```

[38]: from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import LSTM
      from tensorflow.keras.layers import Dense,Input

```

2.2 Initializing the model

```
[39]: model = Sequential()
```

2.3 Adding LSTM Layers

```
[40]: model.add(LSTM(50, return_sequences=True, input_shape=(10,1)))  
model.add(LSTM(50, return_sequences=True))  
model.add(LSTM(50))
```

2.4 Adding output Layers

```
[41]: model.add(Dense(1))
```

```
[42]: model.summary()
```

Model: "sequential_1"

| Layer (type) | Output Shape | Param # |
|-----------------|----------------|---------|
| lstm_3 (LSTM) | (None, 10, 50) | 10,400 |
| lstm_4 (LSTM) | (None, 10, 50) | 20,200 |
| lstm_5 (LSTM) | (None, 50) | 20,200 |
| dense_1 (Dense) | (None, 1) | 51 |

Total params: 50,851 (198.64 KB)

Trainable params: 50,851 (198.64 KB)

Non-trainable params: 0 (0.00 B)

2.5 Configure The Learning Process

```
[43]: model.compile(loss='mean_squared_error', optimizer='adam')
```

2.6 Train The model

```
[44]: model.fit(X_train, y_train, validation_data=(X_test,ytest), epochs=10,  
↳batch_size=64, verbose=1)
```

```
Epoch 1/10  
84/84          12s 43ms/step -  
loss: 0.0055 - val_loss: 8.8414e-04  
Epoch 2/10  
84/84          3s 30ms/step -  
loss: 1.3258e-04 - val_loss: 7.6096e-04  
Epoch 3/10  
84/84          2s 28ms/step -  
loss: 1.1566e-04 - val_loss: 8.1974e-04  
Epoch 4/10  
84/84          3s 31ms/step -  
loss: 1.1804e-04 - val_loss: 7.6955e-04  
Epoch 5/10  
84/84          3s 29ms/step -  
loss: 1.1955e-04 - val_loss: 0.0011  
Epoch 6/10  
84/84          2s 29ms/step -  
loss: 1.1475e-04 - val_loss: 7.9172e-04  
Epoch 7/10  
84/84          2s 28ms/step -  
loss: 1.3041e-04 - val_loss: 8.3297e-04  
Epoch 8/10  
84/84          2s 28ms/step -  
loss: 1.1645e-04 - val_loss: 7.7667e-04  
Epoch 9/10  
84/84          3s 29ms/step -  
loss: 1.1570e-04 - val_loss: 9.2034e-04  
Epoch 10/10  
84/84          3s 32ms/step -  
loss: 1.0593e-04 - val_loss: 0.0027
```

```
[44]: <keras.src.callbacks.history.History at 0x1aaadb4f590>
```

```
[45]: train_predict = model.predict(X_train)  
test_predict = model.predict(X_test)
```

```
167/167          3s 15ms/step  
90/90           1s 8ms/step
```

2.7 Model Evaluation

```
[47]: train_predict=scaler.inverse_transform(train_predict)
      test_predict=scaler.inverse_transform(test_predict)
```

```
[48]: import math
      from sklearn.metrics import mean_squared_error
      math.sqrt(mean_squared_error(y_train,train_predict))
```

```
[48]: 3827.0821056981763
```


2.8 Save the Model

```
[49]: from tensorflow.keras.models import load_model
```

```
[52]: model.save('Crude_oil_price_prediction.keras')
```

2.9 Test the Model

```
[55]: look_back=10
      trainPredictPlot = np.empty_like(data_oil)
      trainPredictPlot[:, :] = np.nan
      trainPredictPlot[look_back: len(train_predict)+look_back, :] = train_predict
      testPredictPlot = np.empty_like(data_oil)
      testPredictPlot[:, :] = np.nan
      testPredictPlot [len (train_predict)+(look_back*2)+1:len(data_oil)-1, :] =
      ↪test_predict
      plt.plot(scaler.inverse_transform(data_oil))
      plt.plot(trainPredictPlot)
      plt.plot(testPredictPlot)
      plt.show()
```

```
[56]: len(test_data)
```

```
[56]: 2876
```

```
[57]: x_input = test_data[2866:].reshape(1,-1)
      x_input.shape
```

```
[57]: (1, 10)
```

```
[58]: temp_input = list(x_input)
      temp_input = temp_input[0].tolist()
```

```
[59]: temp_input
```

```
[59]: [0.44172960165852215,
      0.48111950244335855,
      0.49726047682511476,
      0.4679401747371539,
      0.4729749740855915,
      0.47119798608026064,
      0.47341922108692425,
      0.4649785280616022,
```

```
0.4703835332444839,  
0.47149415074781587]
```

```
[61]: first_output = []  
n_steps = 10  
i = 0  
  
while i < 10:  
    if len(temp_input) > 10:  
        x_input = np.array(temp_input[1:])  
        print(f"{i} day input {x_input}")  
        x_input = x_input.reshape(1, -1)  
        x_input = x_input.reshape((1, n_steps, 1))  
  
        yhat = model.predict(x_input, verbose=0)  
        print(f"{i} day output {yhat}")  
  
        temp_input.extend(yhat[0].tolist())  
        temp_input = temp_input[1:]  
  
        first_output.extend(yhat.tolist())  
        i += 1  
    else:  
        x_input = np.array(temp_input).reshape((1, n_steps, 1))  
        yhat = model.predict(x_input, verbose=0)  
        print(yhat[0])  
  
        temp_input.extend(yhat[0].tolist())  
        print(len(temp_input))  
        first_output.extend(yhat.tolist())  
        i += 1
```

```
[0.43688264]
```

```
11
```

```
1 day input [0.4811195  0.49726048 0.46794017 0.47297497 0.47119799 0.47341922  
0.46497853 0.47038353 0.47149415 0.43688264]  
1 day output [[0.43722245]]  
2 day input [0.49726048 0.46794017 0.47297497 0.47119799 0.47341922 0.46497853  
0.47038353 0.47149415 0.43688264 0.43722245]  
2 day output [[0.4339377]]  
3 day input [0.46794017 0.47297497 0.47119799 0.47341922 0.46497853 0.47038353  
0.47149415 0.43688264 0.43722245 0.4339377 ]  
3 day output [[0.42893893]]  
4 day input [0.47297497 0.47119799 0.47341922 0.46497853 0.47038353 0.47149415  
0.43688264 0.43722245 0.4339377 0.42893893]  
4 day output [[0.42501116]]  
5 day input [0.47119799 0.47341922 0.46497853 0.47038353 0.47149415 0.43688264  
0.43722245 0.4339377 0.42893893 0.42501116]
```

```

5 day output [[0.42015254]]
6 day input [0.47341922 0.46497853 0.47038353 0.47149415 0.43688264 0.43722245
0.4339377 0.42893893 0.42501116 0.42015254]
6 day output [[0.41498962]]
7 day input [0.46497853 0.47038353 0.47149415 0.43688264 0.43722245 0.4339377
0.42893893 0.42501116 0.42015254 0.41498962]
7 day output [[0.4094336]]
8 day input [0.47038353 0.47149415 0.43688264 0.43722245 0.4339377 0.42893893
0.42501116 0.42015254 0.41498962 0.4094336 ]
8 day output [[0.40418002]]
9 day input [0.47149415 0.43688264 0.43722245 0.4339377 0.42893893 0.42501116
0.42015254 0.41498962 0.4094336 0.40418002]
9 day output [[0.39845008]]

```


```
[65]: day_new = np.arange(1, 11)
      day_pred = np.arange(11, 21)
```

```
[66]: len(data_oil)
```

```
[66]: 8216
```


```
[68]: plt.plot(day_new, scaler.inverse_transform(data_oil[8206:]))
      plt.plot(day_pred, scaler.inverse_transform(first_output))
```

```
[68]: [<matplotlib.lines.Line2D at 0x1aab4c23350>]
```



```
[71]: df3 = data_oil.tolist()
      df3.extend(first_output)
      plt.plot(df3[8100:])
```

```
[71]: [<matplotlib.lines.Line2D at 0x1aab4ee61b0>]
```



```
[73]: df3 = scaler.inverse_transform(df3).tolist()
      plt.title('past data had next 10 days output prediction after reversing the_↵
      ↵sacaled values')
      plt.plot(df3)
```

```
[73]: [<matplotlib.lines.Line2D at 0x1aab501e8a0>]
```

past data had next 10 days output prediction after reversing the scaled values

